

# Introduction to machine learning and neural networks

Toby Dylan Hocking  
toby.hocking@nau.edu  
toby.hocking@r-project.org

May 26, 2021

## Introduction and overview

Problem 1: avoiding overfitting in regression



Problem 2: classifying images of digits

Problem 3: predicting earth system model parameters

# Machine learning intro: image classification example

ML is all about learning predictive functions  $f(x) \approx y$ , where

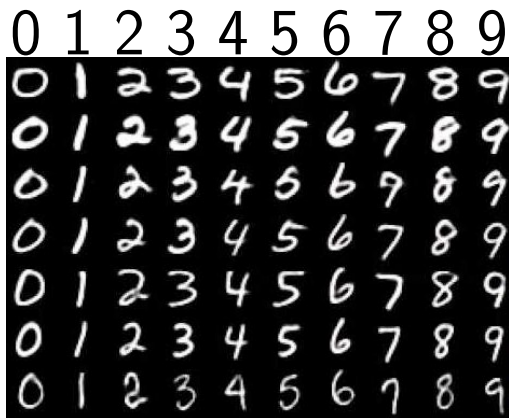
- ▶ Inputs/features  $x$  can be easily computed using traditional algorithms, e.g. matrix of pixel intensities in an image.
- ▶ Outputs/labels  $y$  are what we want to predict, easy to get by asking a human, but hard to compute using traditional algorithms, e.g. image class.
- ▶ Input  $x$  = image of digit, output  $y \in \{0, 1, \dots, 9\}$ ,
  - this is a classification problem with 10 classes.

  $f(\text{image of } 0) = 0$ ,   $f(\text{image of } 1) = 1$

- ▶ Traditional/unsupervised algorithm: I give you a pixel intensity matrix  $x \in \mathbb{R}^{28 \times 28}$ , you code a function  $f$  that returns one of the 10 possible digits. Q: how to do that?

# Supervised machine learning algorithms

I give you a training data set with paired inputs/outputs, e.g.



Your job is to code an algorithm, `LEARN`, that infers a function  $f$  from the training data. (you don't code  $f$ )

Source: [github.com/cazala/mnist](https://github.com/cazala/mnist)

# Advantages of supervised machine learning

Learning Algorithm	Train data	Learned function	Predictions on test data
--------------------	------------	------------------	--------------------------

Learn(



) → g

$g(\text{0}) = 0$

$g(\text{1}) = 1$

$g(\text{1}) = 1$

Learn(



) → h

$h(\text{0}) = 0$

$h(\text{0}) = 0$

$h(\text{1}) = 1$

- ▶ Input  $x \in \mathbb{R}^{28 \times 28}$ , output  $y \in \{0, 1, \dots, 9\}$  types the same!
- ▶ Can use same learning algorithm regardless of pattern.
- ▶ Pattern encoded in the labels (not the algorithm).
- ▶ Useful if there are many un-labeled data, but few labeled data (or getting labels is long/costly).
- ▶ State-of-the-art accuracy (if there is enough training data).

Sources: [github.com/cazala/mnist](https://github.com/cazala/mnist), [github.com/zalando-research/fashion-mnist](https://github.com/zalando-research/fashion-mnist)

# Overview of tutorial

In this tutorial we will discuss two types of problems, which differ by the type of the output/label/ $y$  variable we want to predict.

- ▶ Regression,  $y$  is a real number.
- ▶ Classification,  $y$  is an integer representing a category.

The rest of the tutorial will explain three learning problems:

- ▶ Regression with a single input, to demonstrate how to avoid overfitting.
- ▶ Classification of digit images, to demonstrate how to compare machine learning algorithms in terms of test/prediction accuracy.
- ▶ Regression for predicting earth system model parameters, as a relevant application.

Introduction and overview

Problem 1: avoiding overfitting in regression

Problem 2: classifying images of digits

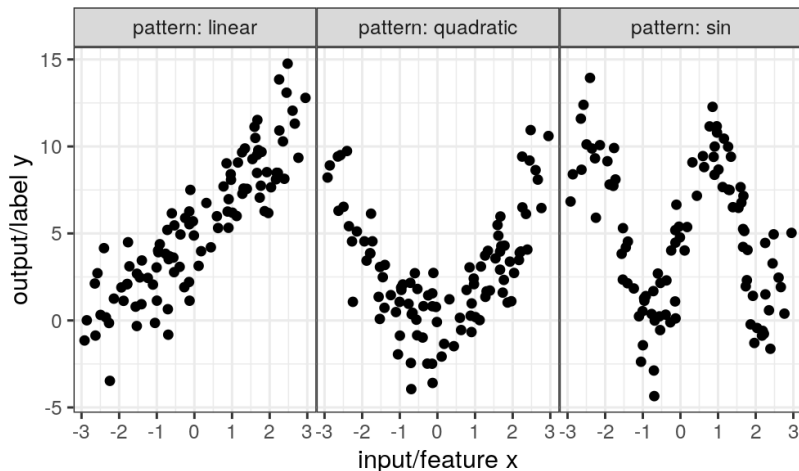
Problem 3: predicting earth system model parameters

# Goal of this section: demonstrate how to avoid overfitting

- ▶ The goal of supervised machine learning is to get accurate predictions on new/unseen/held-out test data.
- ▶ Any machine learning algorithm is prone to overfit, which means providing better predictions on the train/subtrain set than on a held-out validation/test set. (BAD)
- ▶ To learn a model which does NOT overfit (GOOD), you need to first divide your train set into subtrain/validation sets.
- ▶ Code for figures in this section: <https://github.com/tdhock/2020-yiqi-summer-school/blob/master/figure-overfitting.R>



# Three different data sets/patterns



- ▶ We illustrate this using a single input/feature  $x \in \mathbb{R}$ .
- ▶ We use a regression problem with outputs  $y \in \mathbb{R}$ .
- ▶ Goal is to learn a function  $f(x) \in \mathbb{R}$ .

# Neural network prediction function

$$f(\mathbf{x}) = f_L[\cdots f_1[\mathbf{x}]]. \quad (1)$$

With  $L - 1$  hidden layers, we have for all  $l \in \{1, \dots, L\}$ :

$$f_l(t) = A_l(\mathbf{W}_l^\top t), \quad (2)$$


The hyper-parameters which must be fixed prior to learning:

- ▶ Number of layers  $L$ .
- ▶ Activation functions  $A_l$  (classically sigmoid, typically ReLU).
- ▶ Number of hidden units per layer  $(u_1, \dots, u_{L-1})$ .
- ▶ Sparsity pattern in the weight matrices  $\mathbf{W}_l \in \mathbb{R}^{u_l \times u_{l-1}}$ .

The weight matrices  $\mathbf{W}_l$  are learned using gradient descent.


- ▶ In each **iteration** of gradient descent, the weights are updated in order to get better predictions on subtrain data.
- ▶ An **epoch** computes gradients on all subtrain data; there can be from 1 to  $N$ (subtrain size) iterations per epoch.

# Illustration of 4-fold cross-validation



Randomly assign each observation a fold ID from 1 to 4.

# Illustration of subtrain/validation split



- ▶ For validation fold 1, all observations with that fold ID are considered the validation set.
- ▶ All other observations are considered the subtrain set.

# CSV data tables for machine learning

- ▶ One row for each observation.
- ▶ One column for the output/label/ $y$  (in regression the label is a real number, in classification the label is a class/category).
- ▶ The other columns should be inputs/features/ $X$  that will be used to predict the corresponding output/label/ $y$ .

Example: [https://raw.githubusercontent.com/tdhock/2020-yiqi-summer-school/master/data\\_linear.csv](https://raw.githubusercontent.com/tdhock/2020-yiqi-summer-school/master/data_linear.csv)

	x	y
1:	-1.40694802	0.933196336
2:	-0.76725660	3.832773444
3:	0.43712018	4.202983135
4:	2.44924674	13.089055084

...

## Result of reading CSV data into R

Result is:

```
> sim.data
```

	pattern	x	y	fold
1:	linear	-1.4069480	0.9331963	4
2:	linear	-0.7672566	3.8327734	3
3:	linear	0.4371202	4.2029831	1
4:	linear	2.4492467	13.0890551	2
5:	linear	-1.7899084	2.0791987	3
---				
296:	sin	1.7838530	4.0502991	1
297:	sin	-0.2683533	-0.1097264	1
298:	sin	-0.5394955	-0.5539398	1
299:	sin	1.8652215	-0.2262517	4
300:	sin	0.6295997	8.8124249	4

## Assign each observation to subtrain/validation set

```
validation.fold <- 1
sim.data[, set := ifelse(
  fold==validation.fold, "validation", "subtrain")]
> sim.data
```

	pattern	x	y	fold	set
1:	linear	-1.4069480	0.9331963	4	subtrain
2:	linear	-0.7672566	3.8327734	3	subtrain
3:	linear	0.4371202	4.2029831	1	validation
4:	linear	2.4492467	13.0890551	2	subtrain
5:	linear	-1.7899084	2.0791987	3	subtrain
---					
296:	sin	1.7838530	4.0502991	1	validation
297:	sin	-0.2683533	-0.1097264	1	validation
298:	sin	-0.5394955	-0.5539398	1	validation
299:	sin	1.8652215	-0.2262517	4	subtrain
300:	sin	0.6295997	8.8124249	4	subtrain

## Neural network with one hidden layer, 20 hidden units

Use for loops to fit different neural network models for each data set and number of iterations.

```
maxit.values <- 10^seq(0, 4)
pattern.values <- c("linear", "quadratic", "sin")
for(i in maxit.values)for(p in pattern.values){
  pattern.data <- sim.data[pattern==p]
  fit <- nnet::nnet(
    y ~ x,
    pattern.data[set=="subtrain"],
    size=20,      #hidden units
    linout=TRUE,  #for regression
    maxit=i)      #max number of iterations
  ...
}
```



# Neural network prediction function

$$f(\mathbf{x}) = f_L[\cdots f_1[\mathbf{x}]].$$

For the nnet code, we have:

$$f_1(t) = A_1(\mathbf{W}_1^T t),$$

$$f_2(t) = A_2(\mathbf{W}_2^T t),$$

$$f(x) = f_2[f_1(x)] = A_2[\mathbf{W}_2^T A_1(\mathbf{W}_1^T x)].$$

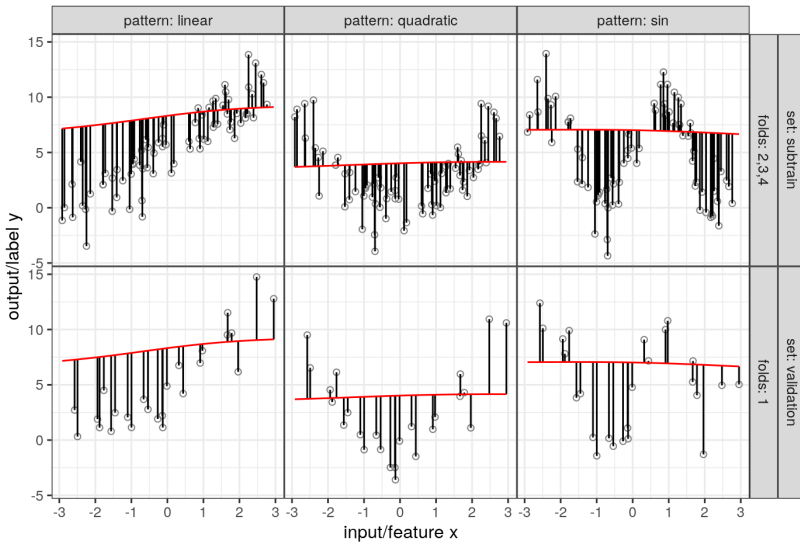
The hyper-parameters are fixed prior to learning:

- ▶ Number of layers  $L = 2$ .
- ▶ Activation functions  $A_1$ =sigmoid,  $A_2$ =identity.
- ▶ Number of units in the hidden layer  $u_1 = 20$ .
- ▶ No sparsity in the weight matrices (fully connected).


The weight matrices  $\mathbf{W}_1, \mathbf{W}_2$  are learned using gradient descent.

- ▶ “Full gradient” method is used, so in each **epoch** there is 1 iteration/update to the weights that is based on the gradient summed over all subtrain data.


# Neural network, 20 hidden units, 1 gradient descent iterations




Neural network, 20 hidden units, 10 gradient descent iterations




## Neural network, 20 hidden units, 100 gradient descent iterations




# Neural network, 20 hidden units, 1000 gradient descent iterations


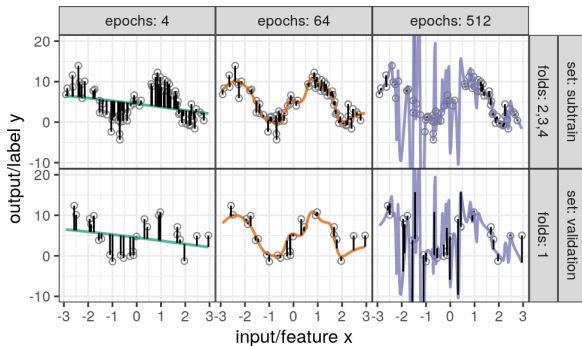


# Neural network, 20 hidden units, 10000 gradient descent iterations



## Neural network, 20 hidden units







# Summary of how to avoid overfitting

- ▶ Happens when subtrain error/loss decreases but validation error increases (as a function of some hyper-parameter)
- ▶ Here the hyper-parameter is the number of iterations of gradient descent, and overfitting starts after a certain number of iterations.
- ▶ To maximize prediction accuracy you need to choose a hyper-parameter with minimal validation error/loss.
- ▶ This optimal hyper-parameter will depend on the data set.
- ▶ To get optimal prediction accuracy in any machine learning analysis, you always need to do this, because you never know the best hyper-parameters in advance.

Introduction and overview

Problem 1: avoiding overfitting in regression

Problem 2: classifying images of digits


Problem 3: predicting earth system model parameters

# Image classification

- ▶ One of the most popular/successful applications of machine learning.
- ▶ Input: image file  $x \in \mathbb{R}^{h \times w \times c}$  where  $h$  is the height in pixels,  $w$  is the width,  $c$  is the number of channels, e.g. RGB image  $c = 3$  channels.
- ▶ In this tutorial we use images with  $h = w = 16$  pixels and  $c = 1$  channel (grayscale, smaller values are darker).
- ▶ Output: class/category  $y$  (from a finite set).
- ▶ In this tutorial there are ten image classes  $y \in \{0, 1, \dots, 9\}$ , one for each digit.



- ▶ Want to learn  $f$  such that  $f(\text{0}) = 0$ ,  $f(\text{1}) = 1$ , etc.
- ▶ Code for figures in this section: <https://github.com/tdhock/2020-yiqi-summer-school/blob/master/figure-validation-loss.R>



# Representation of digits in CSV

- ▶ Each image/observation is one row.
- ▶ First column is output/label/class to predict.
- ▶ Other 256 columns are inputs/features (pixel intensity values).

Data from

<https://web.stanford.edu/~hastie/ElemStatLearn/datasets/zip.train.gz>

```
1:  6 -1 -1  ... -1.000 -1.000  -1
2:  5 -1 -1  ... -0.671 -0.828  -1
3:  4 -1 -1  ... -1.000 -1.000  -1
4:  7 -1 -1  ... -1.000 -1.000  -1
5:  3 -1 -1  ... -0.883 -1.000  -1
6:  6 -1 -1  ... -1.000 -1.000  -1
...
```

# Converting label column to matrix for neural network

This is a “one hot” encoding of the class labels.

```
zip.dt <- data.table::fread("zip.gz")  
zip.y.mat <- keras::to_categorical(zip.dt$V1)
```

	0	1	2	3	4	5	6	7	8	9
[1,]	0	0	0	0	0	0	1	0	0	0
[2,]	0	0	0	0	0	1	0	0	0	0
[3,]	0	0	0	0	1	0	0	0	0	0
[4,]	0	0	0	0	0	0	0	1	0	0
[5,]	0	0	0	1	0	0	0	0	0	0
[6,]	0	0	0	0	0	0	1	0	0	0
...										

# Conversion to array for input to neural network

Use array function with all columns except first as data.

```
zip.size <- 16
zip.X.array <- array(
  data = unlist(zip.dt[1:nrow(zip.dt),-1]),
  dim = c(nrow(zip.dt), zip.size, zip.size, 1))
```

Need to specify dimensions of array:

- ▶ Observations: same as the number of rows in the CSV table.
- ▶ Pixels wide: 16.
- ▶ Pixels high: 16.
- ▶ Channels: 1 (greyscale image).

## Linear model R code

```
library(keras)
linear.model <- keras::keras_model_sequential() %>%
  keras::layer_flatten(
    input_shape = c(16, 16, 1)) %>%
  keras::layer_dense(
    units = 10,
    activation = 'softmax')
```

- ▶ First layer must specify shape of inputs (here 16x16x1).
- ▶ `layer_flatten` converts any shape to a single dimension of units (here 256).
- ▶ `layer_dense` uses all units in the previous layer to predict each unit in the layer.
- ▶ `units=10` because there are ten possible classes for an output.
- ▶ `activation='softmax'` is required for the last/output layer in multi-class classification problems.



## Keras model compilation

```
linear.model %>% keras::compile(  
  loss = keras::loss_categorical_crossentropy,  
  optimizer = keras::optimizer_adadelta(),  
  metrics = c('accuracy')  
)
```

In compile you can specify

- ▶ a loss function, which is directly optimized/minimized in each iteration of the gradient descent learning algorithm.  
<https://keras.io/api/losses/>
- ▶ an optimizer, which is the version of gradient descent learning algorithm to use.  
<https://keras.io/api/optimizers/>
- ▶ an evaluation metric to monitor, not directly optimized via gradient descent, but usually more relevant/interpretable for the application (e.g. accuracy is the proportion of correctly predicted labels). <https://keras.io/api/metrics/>


# Keras model fitting

```
linear.model %>% keras::fit(  
  zip.X.array, zip.y.mat,  
  epochs = 50,  
  validation_split = 0.2  
)
```

In fit you can specify

- ▶ Train data inputs zip.X.array and outputs zip.y.mat (required).
- ▶ Number of full passes of gradient descent through the subtrain data (epochs). In each epoch the gradient with respect to each subtrain observation is computed once.
- ▶ validation\_split=0.2 which means to use 80% subtrain (used for gradient descent parameter updates), 20% validation (used for hyper-parameter selection).

## Linear model




## Sparse (convolutional) model R code

```
library(keras)
conv.model <- keras_model_sequential() %>%
  layer_conv_2d(
    input_shape = dim(zip.X.array)[-1],
    filters = 20,
    kernel_size = c(3,3),
    activation = 'relu') %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_flatten() %>%
  layer_dense(units = 100, activation = 'relu') %>%
  layer_dense(
    units = ncol(zip.y.mat),
    activation = 'softmax')
```

- ▶ Sparse: few inputs are used to predict each unit in layer\_conv\_2d.
- ▶ Exploits structure of image data to make learning easier/faster.


## Convolutional neural network




## Dense (fully connected) neural network R code

```
library(keras)
dense.model <- keras_model_sequential() %>%
  layer_flatten(
    input_shape = dim(zip.X.array)[-1]) %>%
  layer_dense(units = 100, activation = 'relu') %>%
  layer_dense(units = 100, activation = 'relu') %>%
  layer_dense(units = 100, activation = 'relu') %>%
  layer_dense(units = 100, activation = 'relu') %>%
  layer_dense(units = 100, activation = 'relu') %>%
  layer_dense(units = 100, activation = 'relu') %>%
  layer_dense(
    units = ncol(zip.y.mat),
    activation = 'softmax')
```

## Dense (fully connected) neural network with 8 hidden layers








## 4-fold cross-validation for model evaluation

Does the convolutional model provide more accurate predictions on unseen test data? First randomly assign a fold ID to each image/observation, then for each test fold ID from 1 to 4:

- ▶ Hold out the images/observations with the test fold ID as a test set (not used at all for model training).
- ▶ Use the other images/observations to train all three models using `validation_split=0.2`.
- ▶ Plot the validation loss curve as a function of the number of epochs, and select the number of epochs which minimizes the validation loss (hyper-parameter learning).
- ▶ Re-train with `epochs=the learned number of epochs` and `validation_split=0` (all train data used for gradient descent parameter updates).
- ▶ Finally compute the prediction accuracy with respect to the held-out test set.


Plot accuracy values (or mean/sd) after learning/testing four times (one for each test fold).

## Accuracy rates for each test fold



- ▶ Always a good idea to compare with the trivial baseline model which always predicts the most frequent class in the train set. (ignoring all inputs/features)
- ▶ Here we see that the baseline is much less accurate than the three learned models, so they are clearly learning something non-trivial.
- ▶ Code for test accuracy figures: <https://github.com/tdhock/2020-yiqi-summer-school/blob/master/figure-test-accuracy.R>

## Zoom to learned models



- ▶ Dense neural network slightly more accurate than linear model, convolutional significantly more accurate than others.
- ▶ Conclusion: convolutional neural network should be preferred for most accurate predictions in these data.
- ▶ Maybe not the same conclusion in other data sets, with the same models. (always need to do cross-validation experiments to see which model is best in any given data set)
- ▶ Maybe other models/algorithms would be even more accurate in these data. (more/less layers, more/less units, completely different algorithm such as random forests, boosting, etc)

Introduction and overview

Problem 1: avoiding overfitting in regression

Problem 2: classifying images of digits

Problem 3: predicting earth system model parameters

# Problem setting

TODO

# Results

